Aprendizagem 2024

Lab 4: *k*NN and Evaluation

Practical exercises

Consider the following data:

	input		output	
	y_1	y_2	y_3	y_4
x_1	1	1	A	1.4
x_2	2	1	B	0.5
x_3	2	3	B	2
x_4	3	3	B	2.2
x_5	1	0	A	0.7
x_6	1	4	A	1.2

1. Assuming a *k*-nearest neighbor with *k*=3 applied within a leave-one-out schema:

   a) Let y_3 be the output variable (*categoric*). Classify x_1 when considering uniform weights and:

      i. Euclidean (*l*2) distance (real input variables)

      ii. Hamming distance (categorical input variables)

   b) Let y_4 be the output variable (*numeric*). Considering cosine similarity, provide the mean regression estimate for x_1.

   c) Consider a weighted-distance *k*-nearest neighbor with Euclidean (*l*2) distance, identify the:

      i. weighted modeestimate of x_1 for the $y_3$outcome
      ii. weighted mean estimate of x_1 for the $y_4$outcome

2. Let x_j be the measurement on variable y_j for a given observation x.

Given the learnt regression model $\hat{x}_4 = 1 - 0.8x_1 + 0.2{x_2}^2 + 0.2x_1x_2$:

a) Compute the y_4 regression estimates for the observations of the aforementioned dataset

b) Compute the training Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE)

c) Perform a residue analysis to assess the presence of systemic biases against y_1 and y_2

3. Consider the probabilistic outcome of a classifier for the given six observations to be

$$p(y_3 = A \mid x) = \left[p(y_3 = A \mid x_1), \ldots, p(y_3 = A \mid x_6)\right] = [0.45, 0.4, 0.3, 0.6, 0.8, 0.4]$$

a) Draw the training ROC curve

b) Compute the training AUC

c) Would you change the default 0.5 probability threshold for this classifier in order to maximize training F1?

Programming quest

1. Consider the accuracy estimates collected under a 5-fold CV for two predictive models M1 and M2, *acc*M1=(0.7,0.5,0.55,0.55,0.6) and *acc*M2=(0.75,0.6,0.6,0.65,0.55).

Using **scipy**, assess whether the differences in predictive accuracy are statistically significant.

Resource: https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.ttest_rel.html

4. Consider the *housing* dataset available at https://web.ist.utl.pt/~rmch/dscience/data/housing.arff and the *Regression* notebook available at the course's webpage. Using a 10-fold cross-validation:

a) Assess the MAE of a kNN regressor for $k \in \{1,5,9\}$ (remaining parameters as default)

b) Compare the RMSE of the default kNN and decision tree regressors